


# Outlines

- Parameter/Function Optimization
- Learning to Simulate
  - ICLR2019, <https://arxiv.org/abs/1810.02513>
- Simulating, Fast and Slow: Learning Policies for Black-Box Optimization
  - 2024, <https://arxiv.org/abs/2406.04261>
- Black-Box Optimization with Local Generative Surrogates
  - NeurIPS2020, <https://proceedings.neurips.cc/paper/2020/hash/a878dbebc902328b41dbf02aa87abb58-Abstract.html>


# Parameter/Function Optimization



Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- We adjust the parameters  $x$  or  $\theta$  to achieve desired output  $y$ .


# Parameter/Function Optimization



Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- We adjust the parameters  $x$  or  $\theta$  to achieve desired output  $y$ .
  - $x$ : Input to the function,  $\theta$ : Model parameters
    - In some cases,  $x$  and  $\theta$  may be indistinguishable.


# Parameter/Function Optimization



Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- We adjust the parameters  $x$  or  $\theta$  to achieve desired output  $y$ .
  - $x$ : Input to the function,  $\theta$ : Model parameters
    - In some cases,  $x$  and  $\theta$  may be indistinguishable.
  - $y$ : Function output,  $t$ : Desired target


# Parameter/Function Optimization



Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- We adjust the parameters  $x$  or  $\theta$  to achieve desired output  $y$ .
  - $x$ : Input to the function,  $\theta$ : Model parameters
    - In some cases,  $x$  and  $\theta$  may be indistinguishable.
  - $y$ : Function output,  $t$ : Desired target
  - $L$ : Loss or objective function
    - Measures the difference between  $y$  and  $t$


# Parameter/Function Optimization



Ex.  $y = f(x; \theta) = ax^2 + bx + c$


- Parameter Optimization: Focuses on optimizing  $\theta$  to achieve the desired  $y$  for a given  $x$ .
  - Classification: For  $x = \text{🍏}$  and  $t = \text{apple}$ , adjust  $\theta$  so that  $y = \text{apple}$ .

# Parameter/Function Optimization




Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- Function Optimization: Focuses on finding  $x$  such that  $y$  is minimized or maximized.
  - Typical optimization problems.
  - $t$  is implicitly given as  $\pm\infty$ .



# Parameter/Function Optimization




Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- In our case,
  - $f$ : A traffic simulator
  - $y$ : Simulation output
  - $t$ : Real traffic observation
  - We need to define  $x$ ,  $\theta$  and  $L(y, t)$ .



# Parameter/Function Optimization



Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- Technical Challenge
  - Difficulties arise when  $f$  is a black model and non-differentiable with respect to  $x$  or  $\theta$ .


- Minimization of a function.

$$f(x) = x^2 + 3x + 3$$


$$\frac{\partial f}{\partial x} = 2x + 3 = 0$$

$$x = -\frac{3}{2}$$


$$f(x) = x^2 + 3x + 3$$




# Gradient Descent/Ascent



$$x_{n+1} = x_n - \lambda \left. \frac{\partial f}{\partial x} \right|_{x_n}$$



# Parameter/Function Optimization



Ex.  $y = f(x; \theta) = ax^2 + bx + c$

- There are several approaches to optimize a non-differentiable black box function.
  - Policy Gradient/REINFORCE
  - Bayesian Optimization
  - Local Surrogate-based Approach
  - Genetic Algorithm
  - Particle Swam Optimization

# Outlines

- Parameter/Function Optimization
- Learning to Simulate
  - ICLR2019, <https://arxiv.org/abs/1810.02513>
- Simulating, Fast and Slow: Learning Policies for Black-Box Optimization
  - 2024, <https://arxiv.org/abs/2406.04261>
- Black-Box Optimization with Local Generative Surrogates
  - NeurIPS2020, <https://proceedings.neurips.cc/paper/2020/hash/a878dbebc902328b41dbf02aa87abb58-Abstract.html>

# Learning to Simulate

- ICLR2019, <https://arxiv.org/abs/1810.02513>

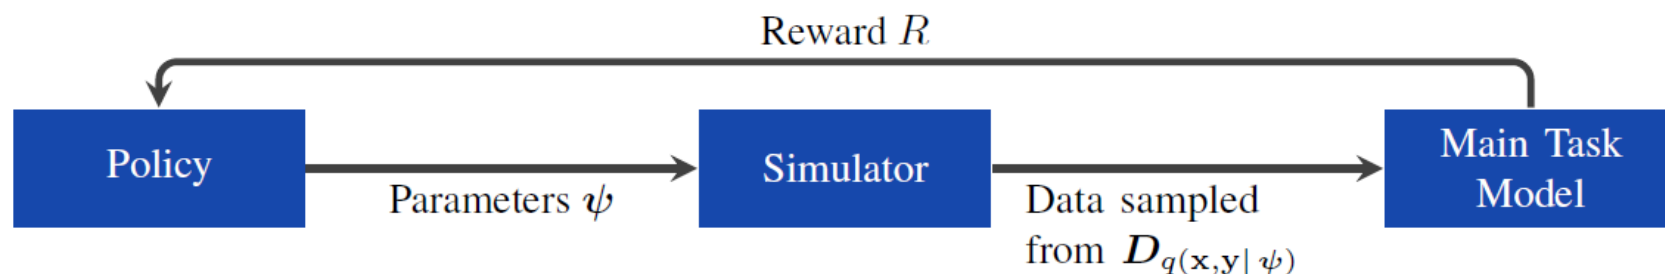


Figure 1: A high-level overview of our “learning to simulate” approach. A policy  $\pi_{\omega}$  outputs parameters  $\psi$  which are used by a simulator to generate a training dataset. The main task model (MTM) is then trained on this dataset and evaluated on a validation set. The obtained accuracy serves as reward signal  $R$  for the policy on how good the synthesized dataset was. The policy thus learns *how to generate data to maximize the validation accuracy*.

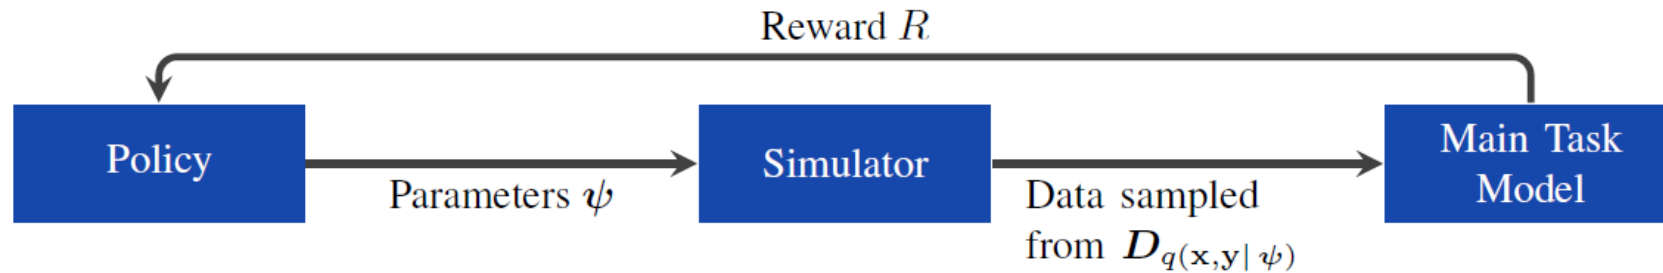
- Simulator: Traffic Scene Generation
  - Car Counting, Segmentation



Figure 3: Example of rendered traffic scene with CARLA (Dosovitskiy et al., 2017) and the Unreal engine (Epic-Games, 2018).

- A straight road of variable length.
- Either an L, T or X intersection at the end of the road.
- Cars of 5 different types which are spawned randomly on the straight road.
- Houses of a unique type which are spawned randomly on the sides of the road.
- Four different types of weather.

# Learning to Simulate



**Policy:** Determines the simulation setting parameters  $\psi$ .

**Simulator:** Generate a dataset based on parameters  $\psi$ .


**Main Task Model:** Trains on the data generated from the simulator for task like:  
Classification, Regression, Segmentation

**Reward  $\psi$ :** The feedback signal is based on the performance of the **Main Task** model (Accuracy/Error).



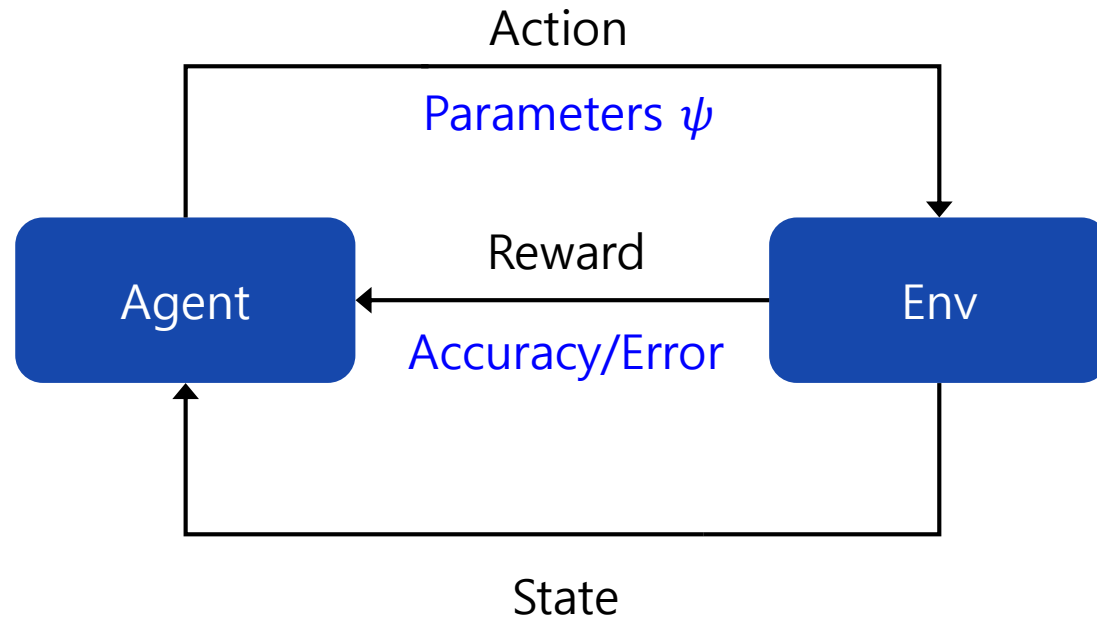
# Learning to Simulate

- RL-based Approach



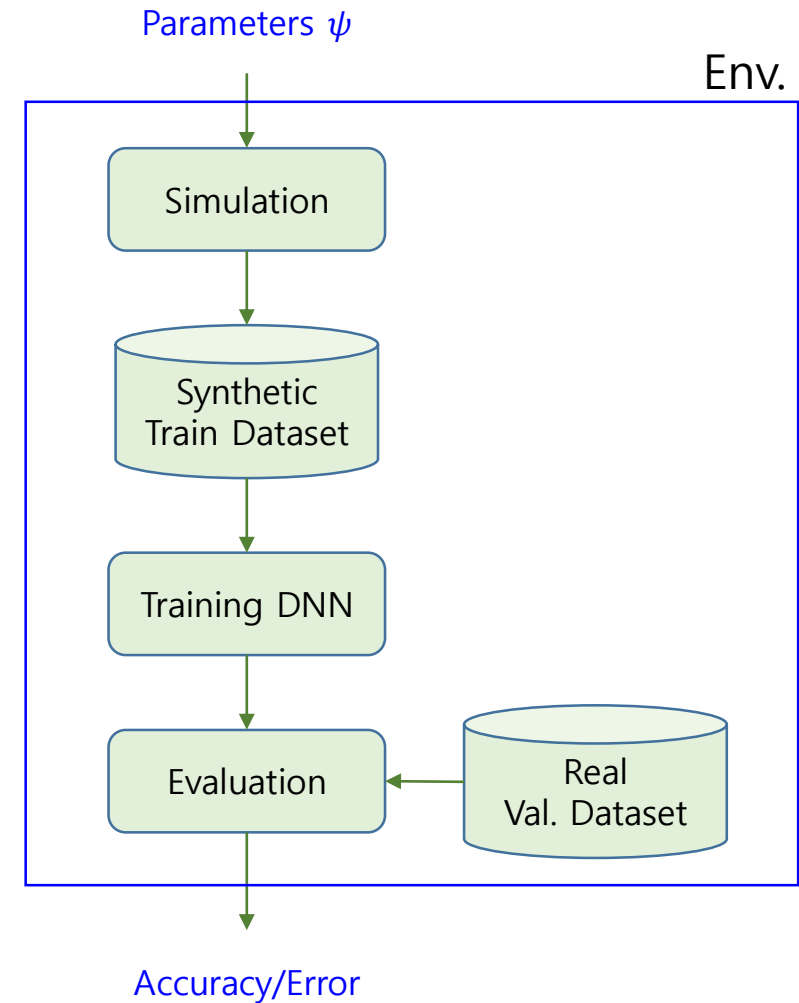
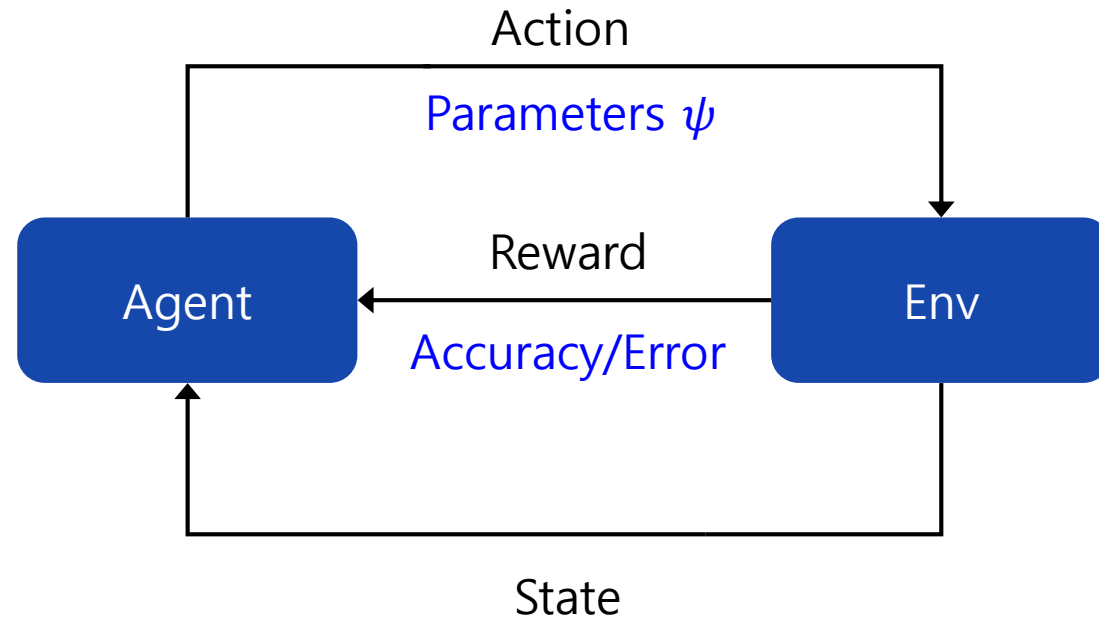
# Learning to Simulate

- RL-based Approach



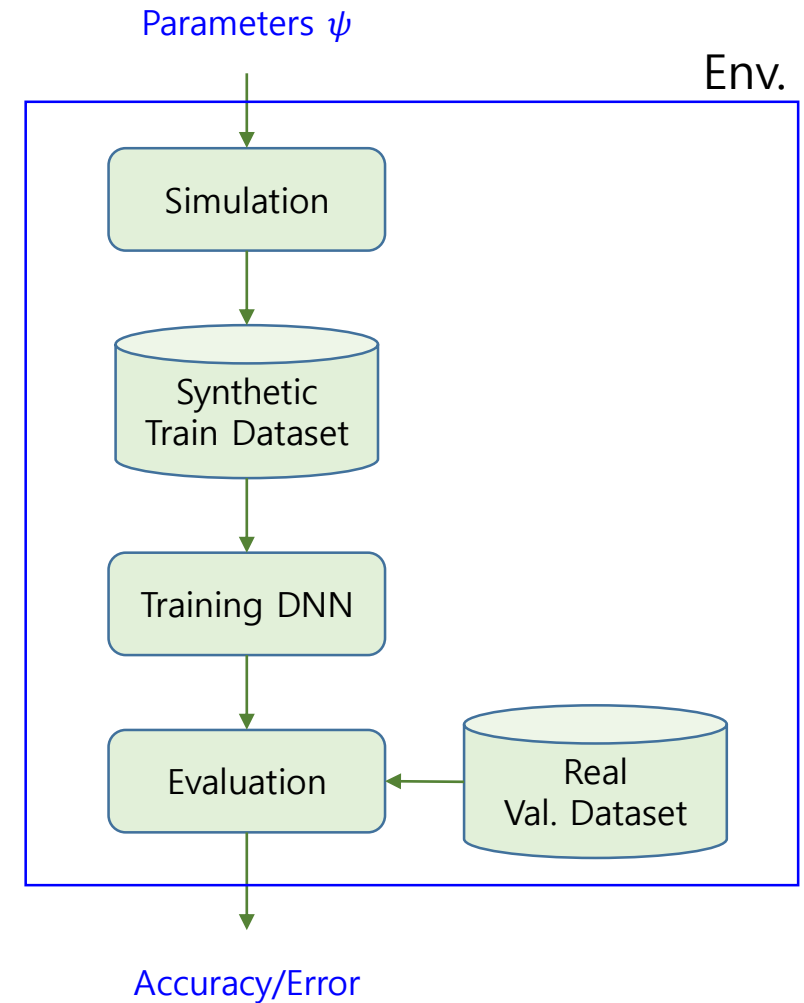
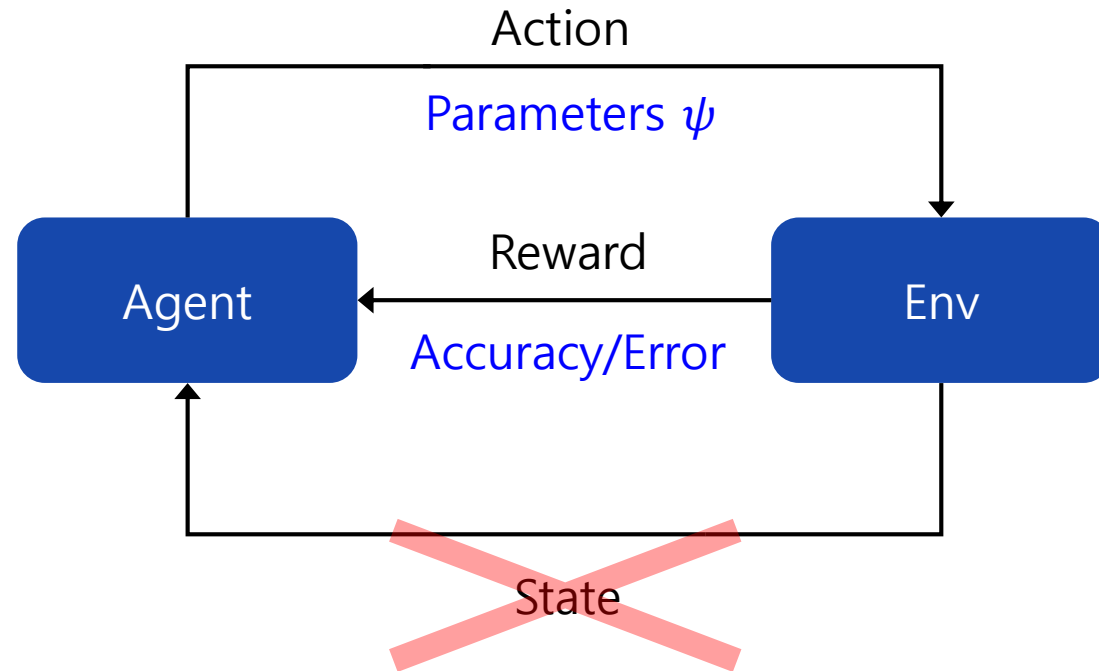
# Learning to Simulate

- RL-based Approach





# Learning to Simulate

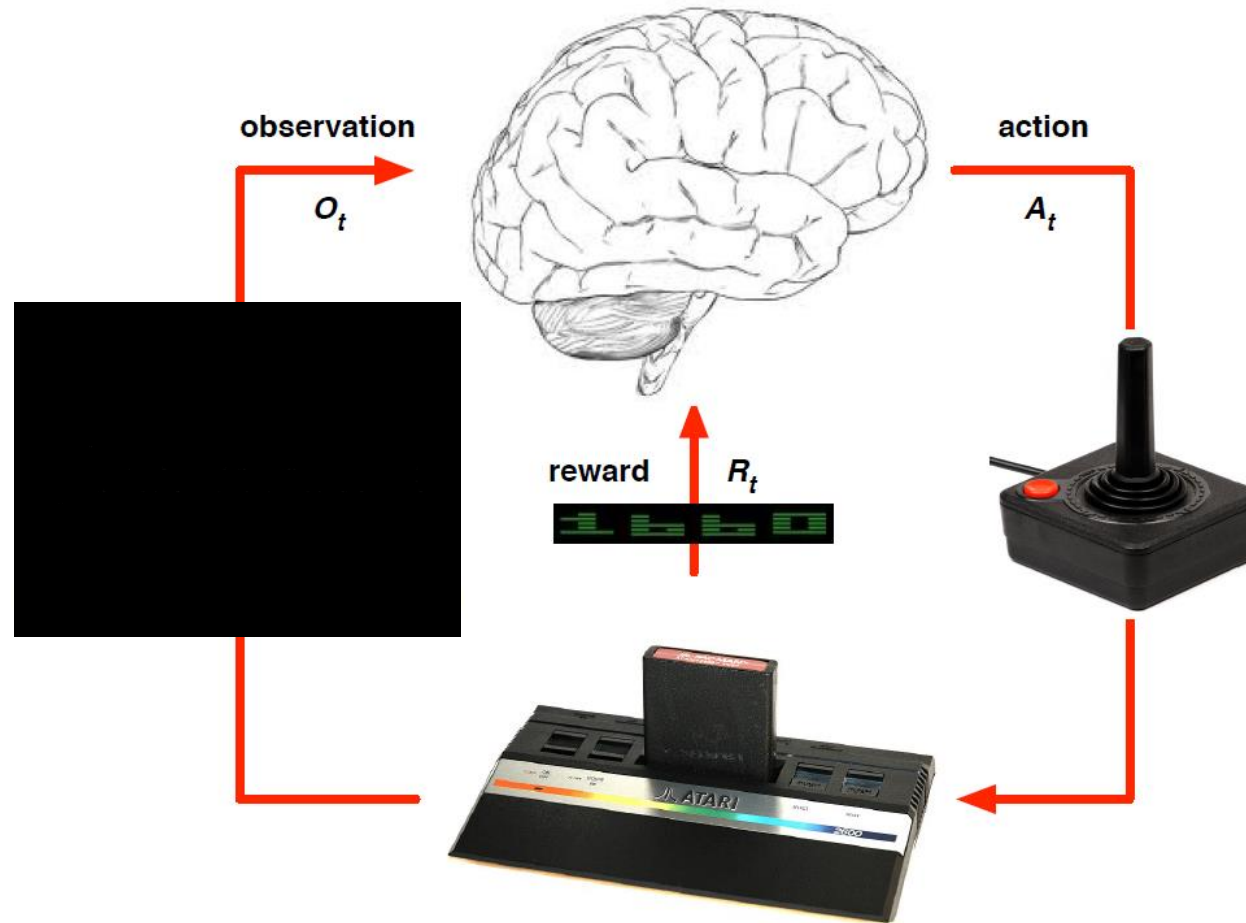
- RL-based Approach



# Learning to Simulate


- RL-based Approach





# Sequential Decision Problem

- Generalizes the previous example.
  - State, Action, Reward.
  - Maximize the sum of rewards.
  - Each decision affects subsequent decisions.



# Learning to Simulate

- Reinforcement Learning (RL) is overkill for this application.
  - Non-sequential problem
    - The simulation optimization problem is **not a sequential decision-making task**.
    - Since the problem doesn't involve sequential decisions, there's no strong reason to use RL.
  - Challenges with RL
    - **Complex training**: Training an RL agent requires significant care to ensure convergence, avoid overfitting, and balance exploration-exploitation.
    - **High computational cost**: RL training often involves extensive simulation calls, making it time-consuming and expensive.
    - **On-policy constraints**: Practical and popular methods like **PPO** require new simulation samples at each iteration, further increasing cost and complexity.

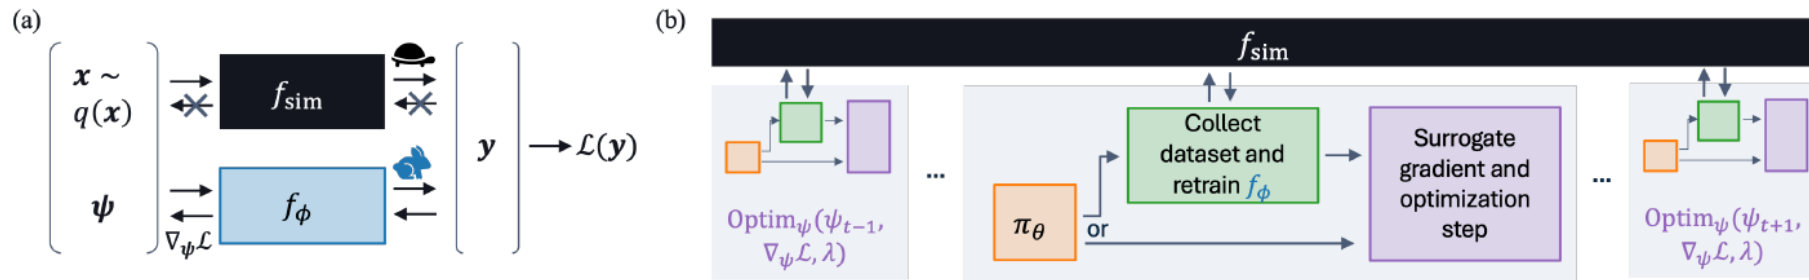


# Outlines

- Parameter/Function Optimization
- Learning to Simulate
  - ICLR2019, <https://arxiv.org/abs/1810.02513>
- Simulating, Fast and Slow: Learning Policies for Black-Box Optimization
  - 2024, <https://arxiv.org/abs/2406.04261>
- Black-Box Optimization with Local Generative Surrogates
  - NeurIPS2020, <https://proceedings.neurips.cc/paper/2020/hash/a878dbebc902328b41dbf02aa87abb58-Abstract.html>

# Simulating, Fast and Slow: Learning Policies for Black-Box Optimization



- 2024, <https://arxiv.org/abs/2406.04261>
  - Surrogate-based Approach



**Figure 1: Schematic view of our approach.** (a) We study black-box optimization problem (over parameters  $\psi$ ), with an emphasis on using gradient information from a fast differentiable surrogate  $f_\phi$  (b) To optimize  $\psi$  sample-efficiently, we employ a policy  $\pi_\theta$  to actively determine whether retraining the surrogate is necessary before using the gradient information.

- Wireless Communication: Indoor Transmitting Antenna Placement

Find the transmitting antenna location ( $\psi$ )  
to maximize signal strength ( $y$ ) at receiver locations ( $x$ ).

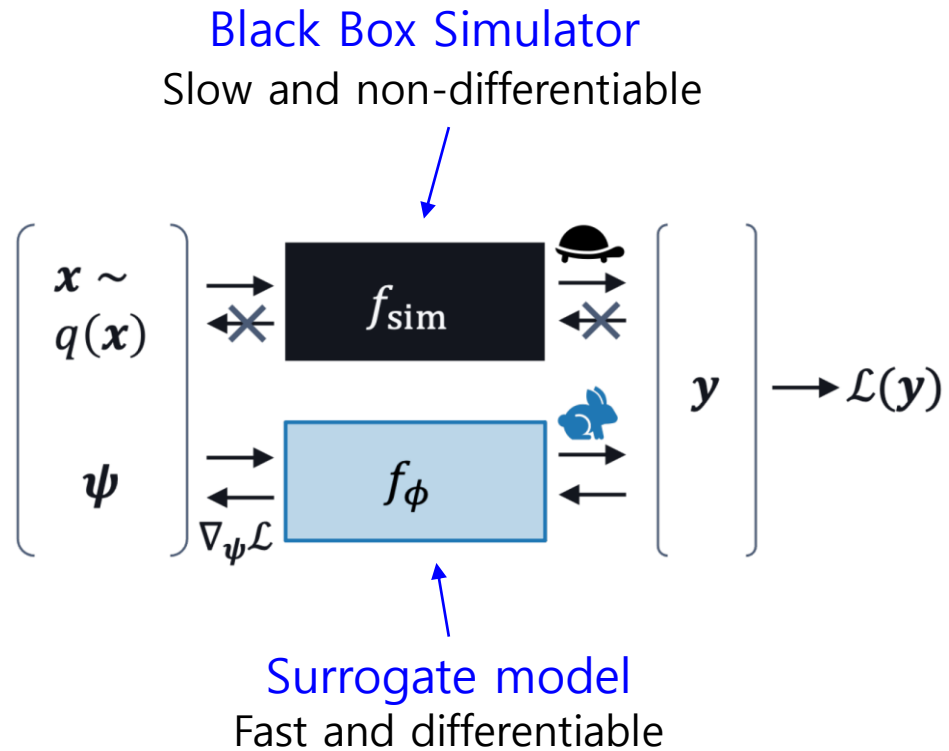


$\psi$ : Transmitting antenna location.

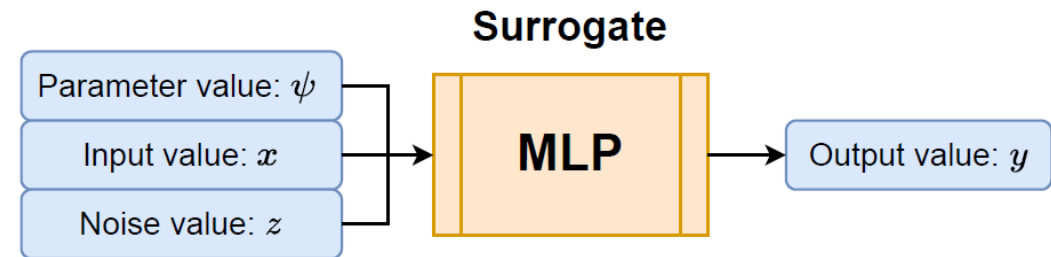
$x$ : Receiving antenna locations.

$y$ : Signal strength at receiver's locations.


# Surrogate Model




- Train a DNN surrogate to approximate the simulator.
- Then, gradients from the surrogate are used for optimizing  $\psi$ .



# Gradient Descent/Ascent

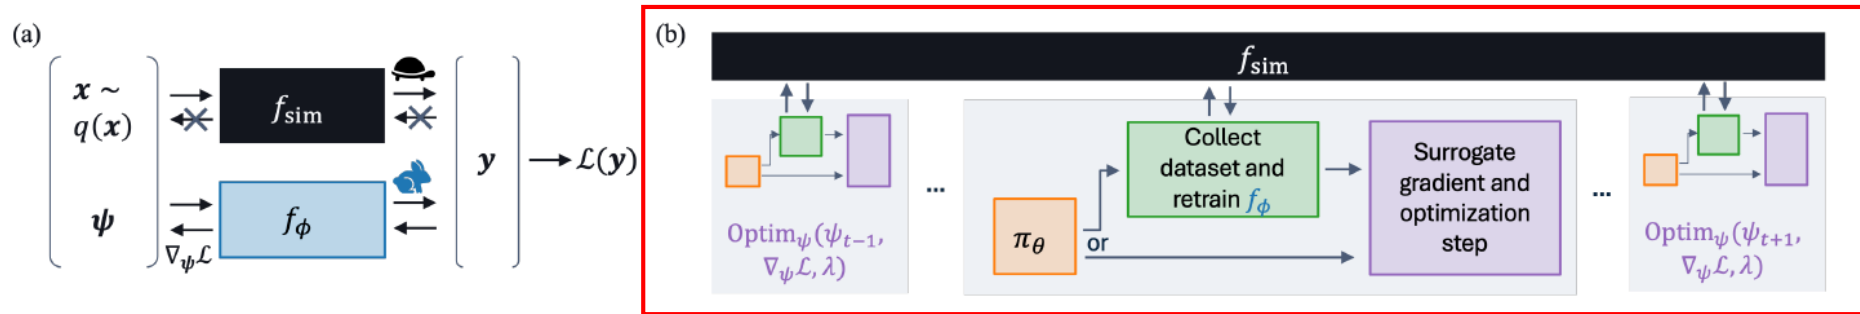


$$\psi_{n+1} = \psi_n - \lambda \left. \frac{\partial f}{\partial \psi} \right|_{\psi_n}$$



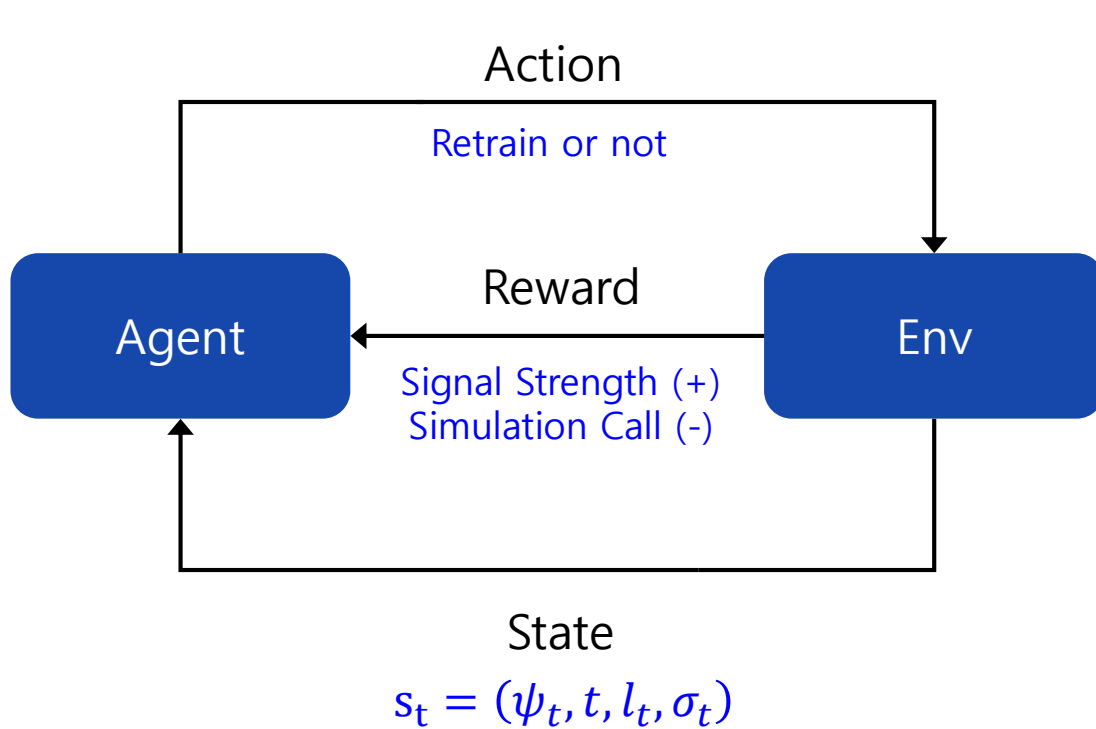
# Learning Policy to train a surrogate model

- Goal: Minimize # simulation calls
  - When to call a simulation
  - How to sample training data.



**Figure 1: Schematic view of our approach.** (a) We study black-box optimization problem (over parameters  $\psi$ ), with an emphasis on using gradient information from a fast differentiable surrogate  $f_\phi$  (b) To optimize  $\psi$  sample-efficiently, we employ a policy  $\pi_\theta$  to actively determine whether retraining the surrogate is necessary before using the gradient information.

# Learning Policy to train a surrogate model

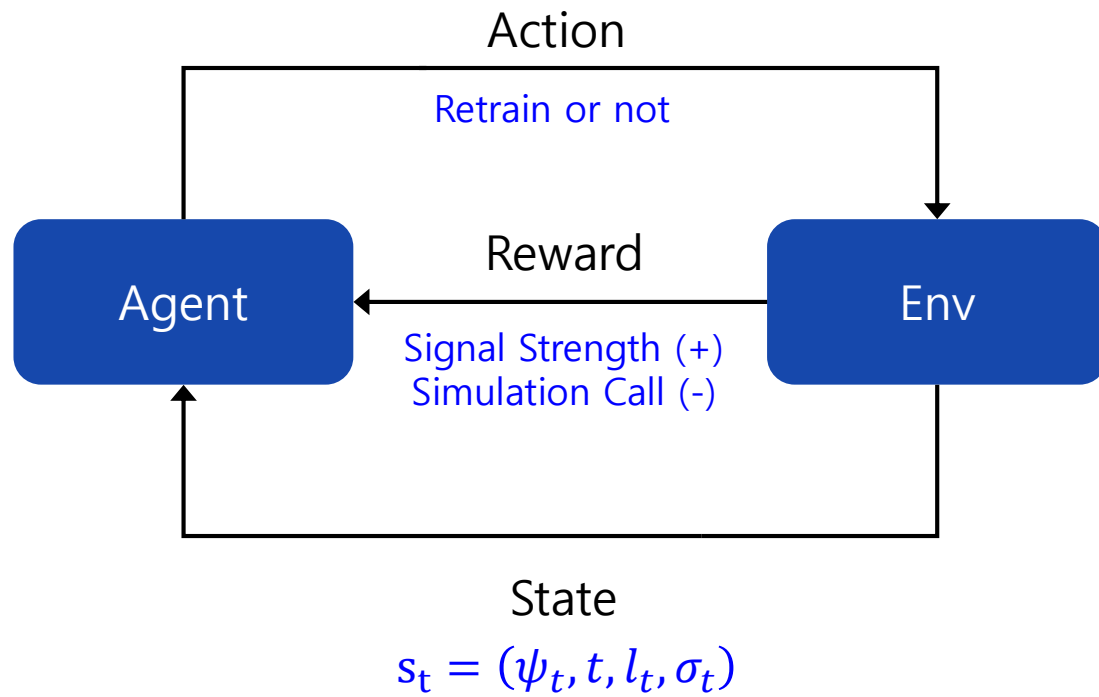


## State

- $\psi_t$ : The current parameter values being optimized.
- $t$ : The current timestep in the optimization process.
- $l_t$ : The number of simulator calls that have already been made during the current episode.
- $\sigma_t$ : A measure of the uncertainty of the surrogate model, indicating how reliable its predictions are at the current  $\psi_t$ .

Note: Agent  $\neq$  Surrogate

## Learning Policy to train a surrogate model



### Action: Retrain

- 1. Determine the sampling boundary for  $\psi_t$ .
- 2. Run simulations
- 3. Collect training data
- 4. Retrain a surrogate model

Note: Agent  $\neq$  Surrogate



# Simulating, Fast and Slow: Learning Policies for Black-Box Optimization

- ~~Reinforcement Learning (RL) is overkill for this application.~~
  - ~~Non-sequential problem~~
    - ~~The simulation optimization problem is **not a sequential decision-making task**.~~
    - ~~Since the problem doesn't involve sequential decisions, there's no strong reason to use RL.~~
- Challenges with RL
  - **Complex training:** Training an RL agent requires significant care to ensure convergence, avoid overfitting, and balance exploration-exploitation.
  - ~~High computational cost:~~ ~~RL training often involves extensive simulation calls, making it time consuming and expensive.~~
  - **On-policy constraints:** Practical and popular methods like **PPO** require new simulation samples at each iteration, further increasing cost and complexity.

# Outlines

- Parameter/Function Optimization
- Learning to Simulate
  - ICLR2019, <https://arxiv.org/abs/1810.02513>
- Simulating, Fast and Slow: Learning Policies for Black-Box Optimization
  - 2024, <https://arxiv.org/abs/2406.04261>
- Black-Box Optimization with Local Generative Surrogates
  - NeurIPS2020, <https://proceedings.neurips.cc/paper/2020/hash/a878dbebc902328b41dbf02aa87abb58-Abstract.html>

# Local Generative Surrogate Optimization

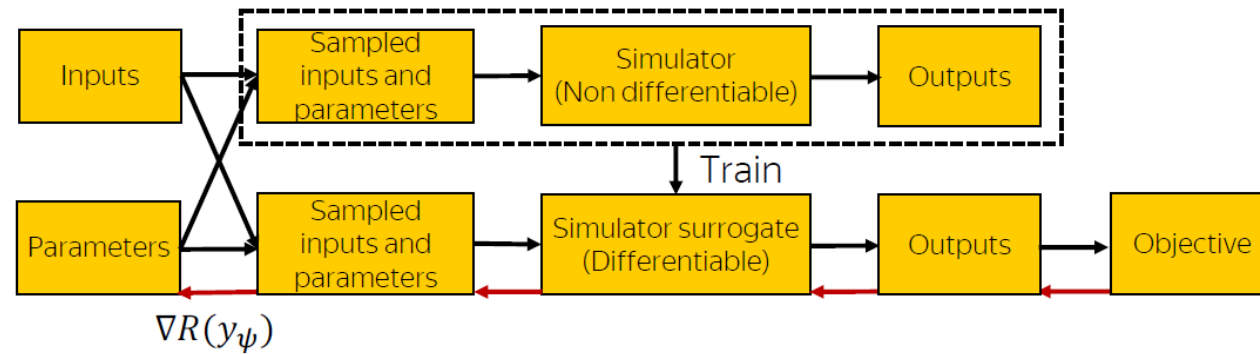
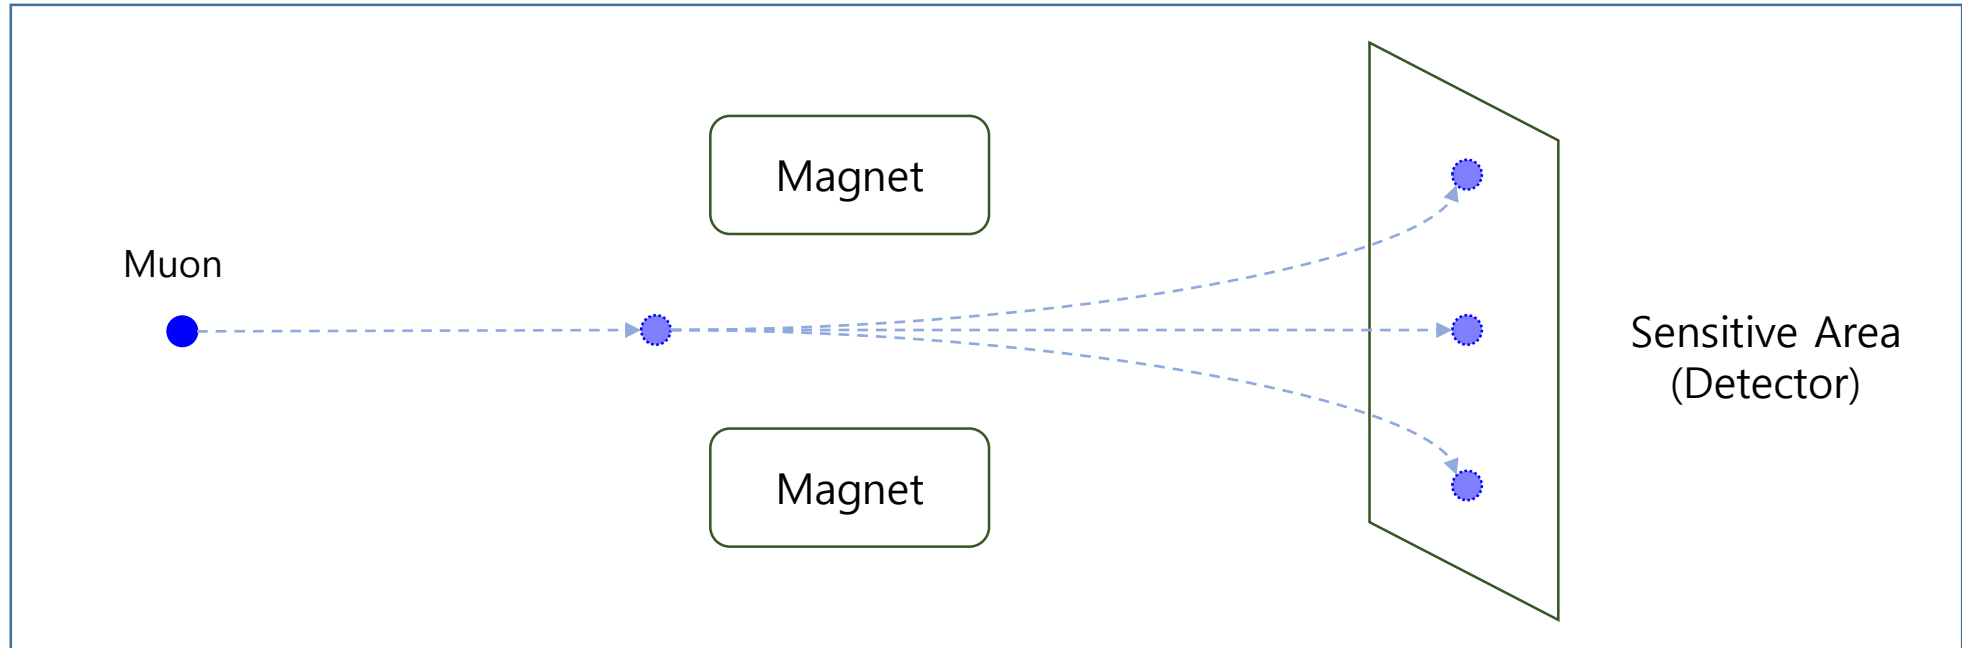


Figure 1: Simulation and surrogate training. *Black*: forward propagation. *Red*: error backpropagation.

# Muon Background Reduction

Minimize # muon hits on sensitive area

$$\mathbf{y} = F(\mathbf{x}; \boldsymbol{\psi})$$



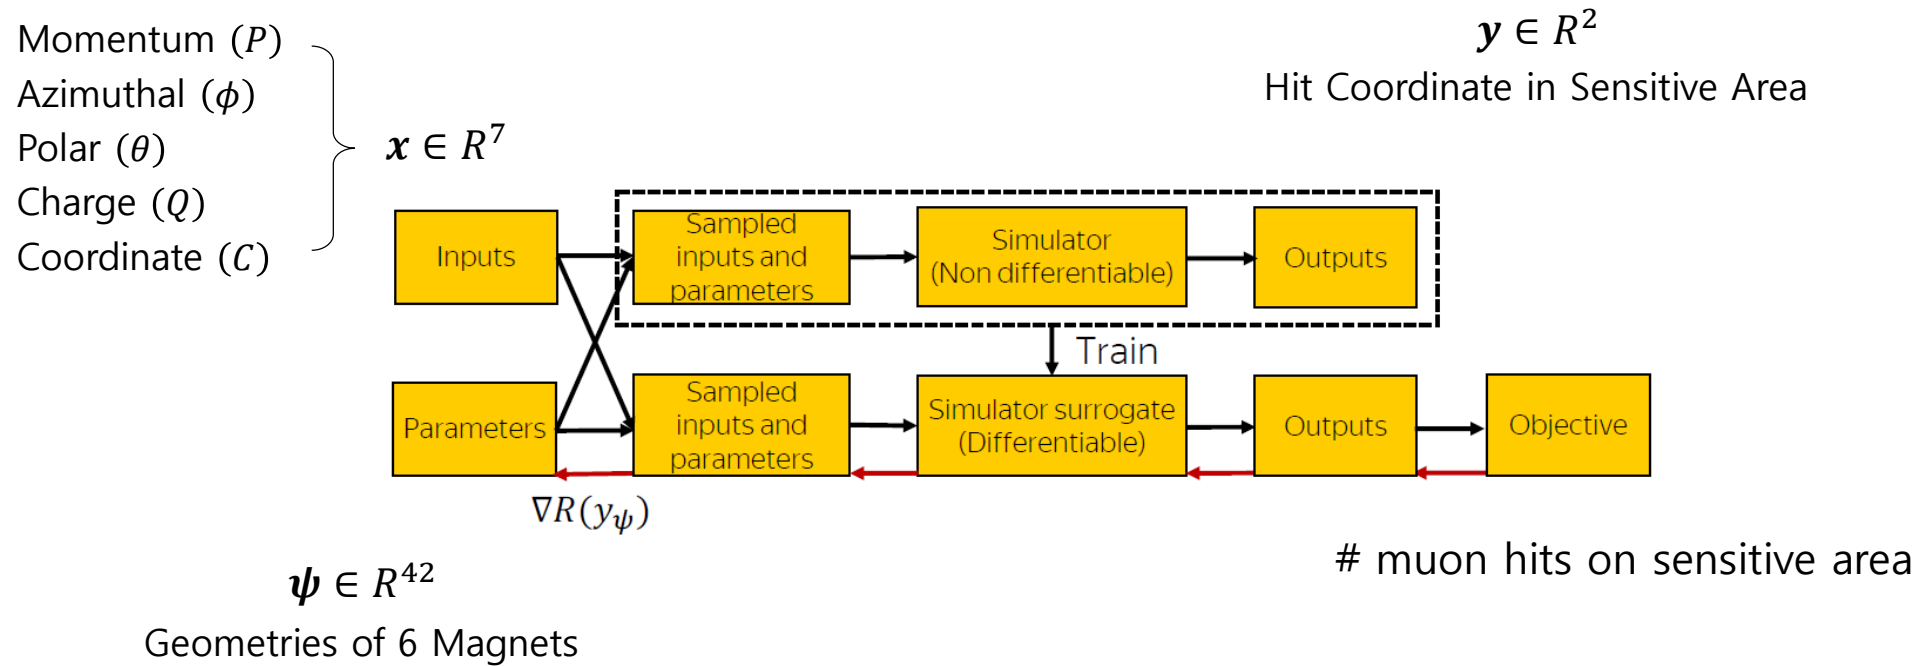
Momentum ( $P$ )  
 Azimuthal ( $\phi$ )  
 Polar ( $\theta$ )  
 Charge ( $Q$ )  
 Coordinate ( $C$ )

$\mathbf{x} \in R^7$

$\boldsymbol{\psi} \in R^{42}$   
 Geometries of 6 Magnets

$\mathbf{y} \in R^2$   
 Hit Coordinate in Sensitive Area

# Local Generative Surrogate Optimization



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 





# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization

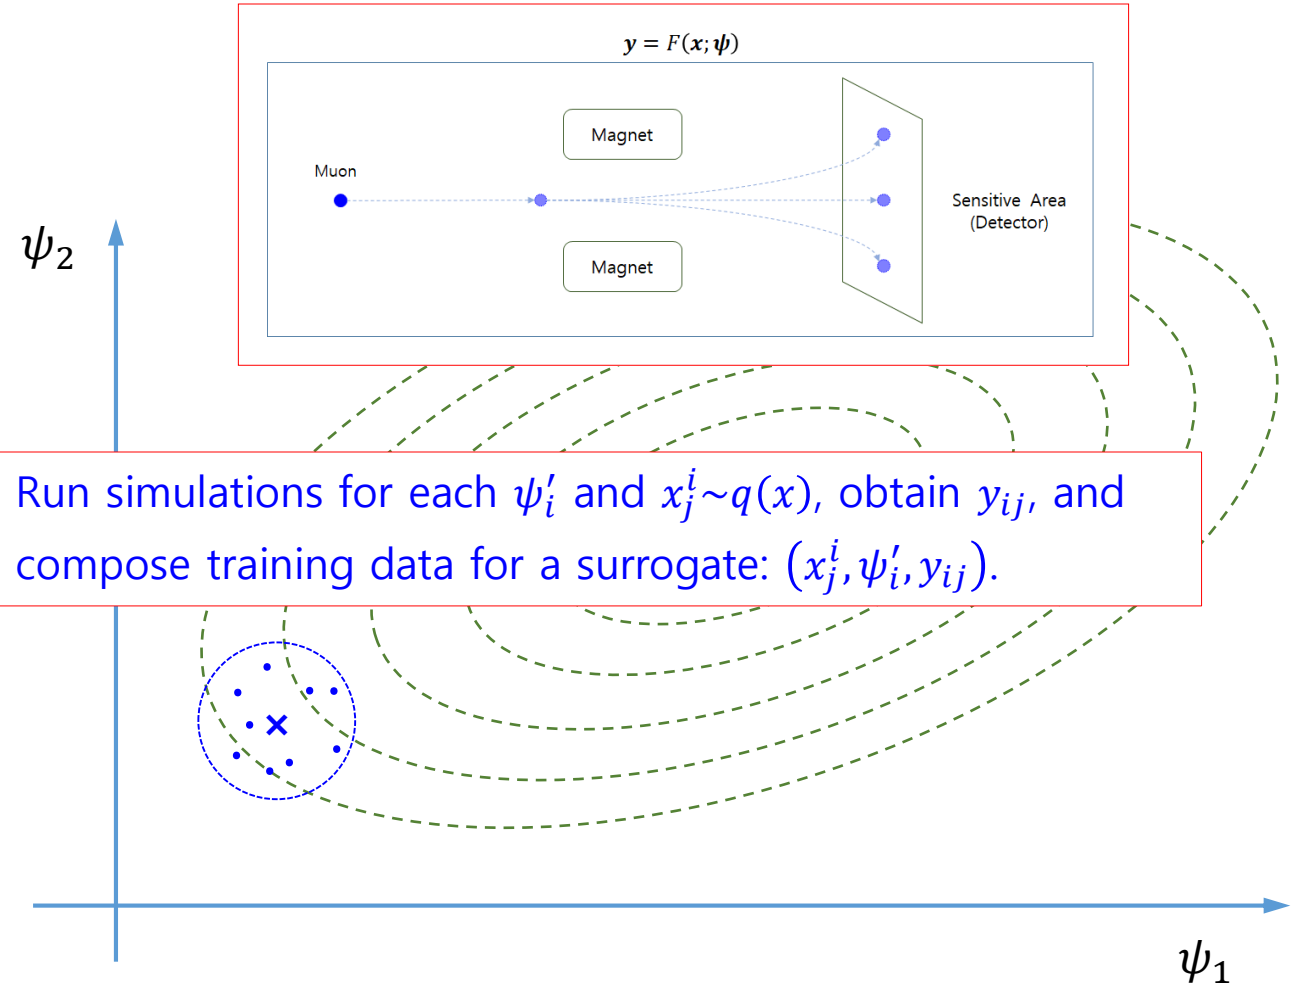
---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization

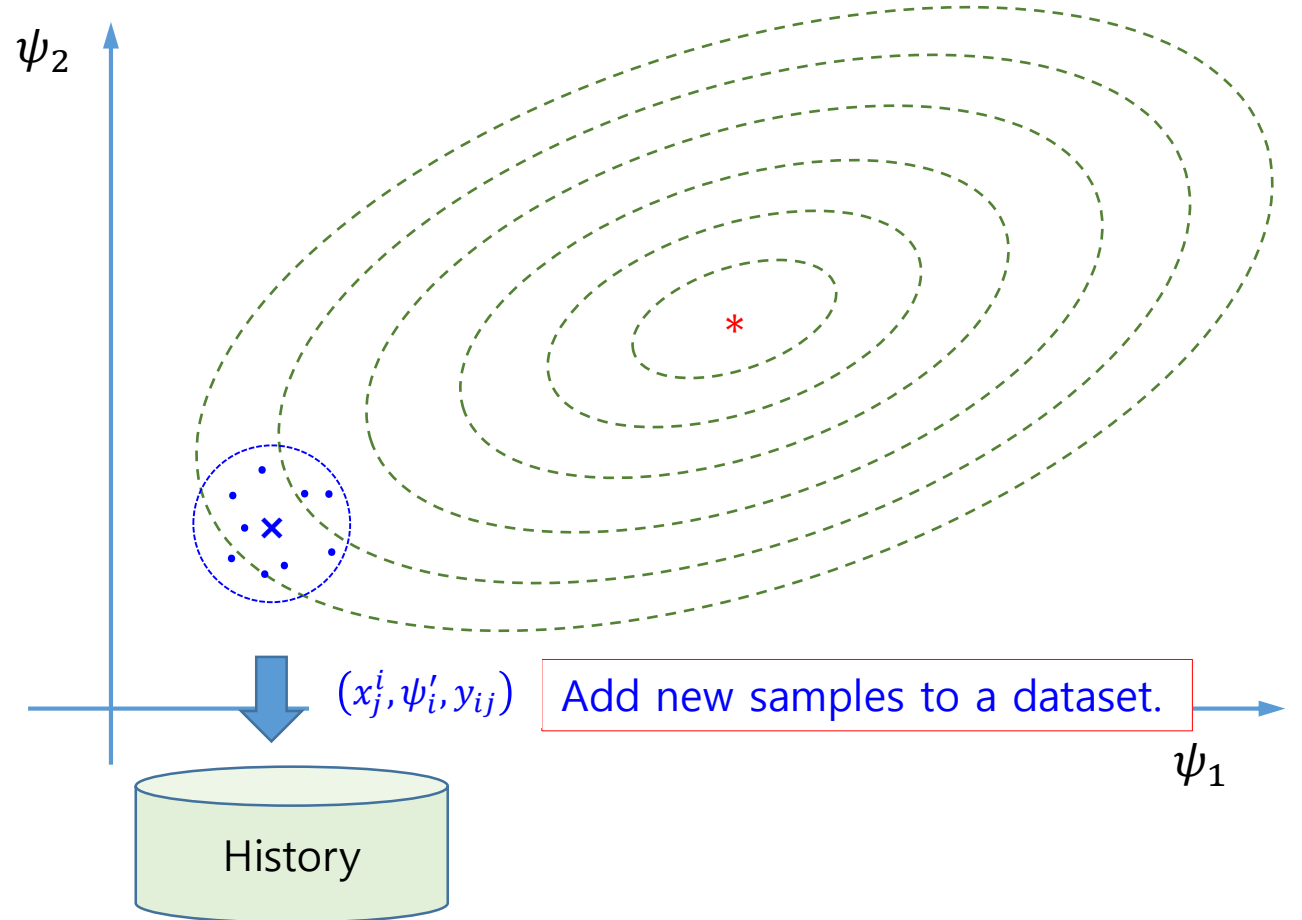
---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization

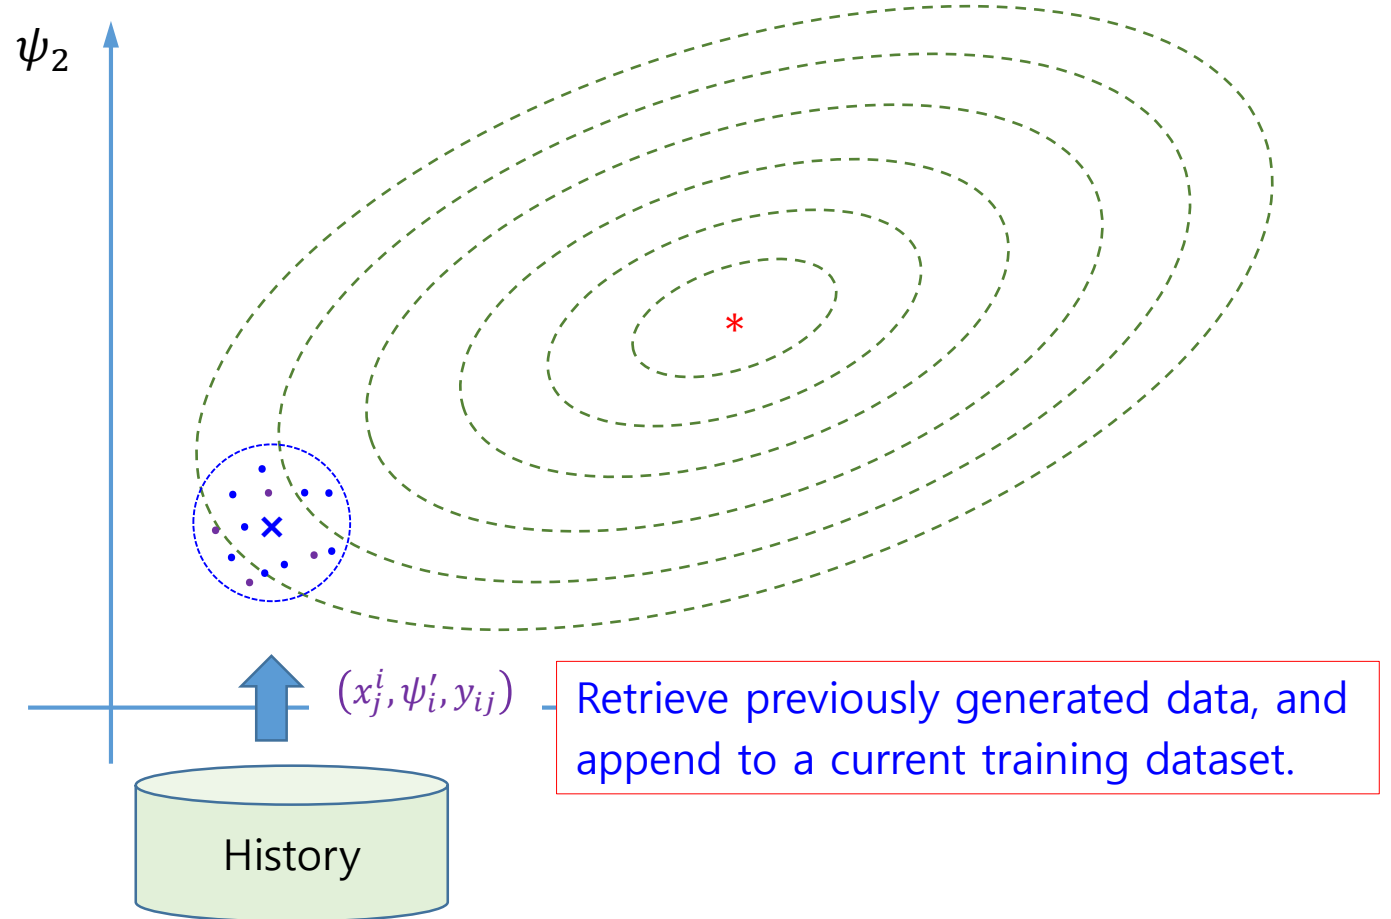
---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



Evaluate a loss function ( $R$ , # muon hits) with  $\bar{\mathbf{y}}$ , and compute the gradient of  $R$  w.r.t  $\psi$ .



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 





# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 



# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 





# Local Generative Surrogate Optimization


---

## Algorithm 1 Local Generative Surrogate Optimization (L-GSO) procedure

---

**Require:** number  $N$  of  $\psi$ , number  $M$  of  $\mathbf{x}$  for surrogate training, number  $K$  of  $\mathbf{x}$  for  $\psi$  optimization step, trust region  $U_\epsilon$ , size of the neighborhood  $\epsilon$ , Euclidean distance  $d$

- 1: Choose initial parameter  $\psi$
  - 2: **while**  $\psi$  has not converged **do**
  - 3:   Sample  $\psi'_i$  in the region  $U_\epsilon^\psi$ ,  $i = 1, \dots, N$
  - 4:   For each  $\psi'_i$ , sample inputs  $\{\mathbf{x}_j^i\}_{j=1}^M \sim q(\mathbf{x})$
  - 5:   Sample  $M \times N$  training examples from simulator  $\mathbf{y}_{ij} = F(\mathbf{x}_j^i; \psi'_i)$
  - 6:   Store  $\mathbf{y}_{ij}, \mathbf{x}_j^i, \psi'_i$  in history  $H$   
 $i = 1, \dots, N; j = 1, \dots, M$
  - 7:   Extract all  $\mathbf{y}_l, \mathbf{x}_l, \psi'_l$  from history  $H$ ,  
 iff  $d(\psi, \psi'_l) < \epsilon$
  - 8:   Train generative surrogate model  
 $S_\theta(\mathbf{z}_l, \mathbf{x}_l; \psi'_l)$ , where  $\mathbf{z}_l \sim \mathcal{N}(0, 1)$
  - 9:   Fix weights of the surrogate model  $\theta$
  - 10:   Sample  $\bar{\mathbf{y}}_k = S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)$ ,  $\mathbf{z}_k \sim \mathcal{N}(0, 1)$ ,  
 $\mathbf{x}_k \sim q(\mathbf{x})$ ,  $k = 1, \dots, K$
  - 11:    $\nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})] \leftarrow \frac{1}{K} \sum_{k=1}^K \frac{\partial \mathcal{R}}{\partial \bar{\mathbf{y}}_k} \frac{\partial S_\theta(\mathbf{z}_k, \mathbf{x}_k; \psi)}{\partial \psi}$
  - 12:    $\psi \leftarrow \text{SGD}(\psi, \nabla_\psi \mathbb{E}[\mathcal{R}(\bar{\mathbf{y}})])$
  - 13: **end while**
- 





Q & A